

COS40007 Artificial Intelligence for Engineering

Portfolio Assessment 1: "Hello Machine Learning for Engineering"

Student Name: Nguyen Quy Hung

Student Number: 104850199

Studio Class: Studio 1&2

Submission Date: September 14, 2025

1. Dataset Selection

Selected Dataset: Water Potability Dataset (water_potability.csv)

I selected a dataset comprising 3,276 water samples with 9 physicochemical parameters and 1 target variable indicating water potability. The original features include: pH, Hardness, Solids, Chloramines, Sulfate, Conductivity, Organic_carbon, Trihalomethanes, and Turbidity.

2. Rationale for Dataset Selection

I chose the Water Potability Dataset because it directly relates to environmental engineering applications, which aligns with my interests in data-driven solutions for real-world problems. As a third-year engineering student, I find water quality assessment particularly relevant since it represents a critical domain where machine learning can provide substantial value in automated monitoring systems. The binary classification problem of determining water safety based on quantitative chemical measurements reflects scenarios I might encounter in my future career, where data-driven approaches can complement traditional analytical methods in environmental monitoring and regulatory compliance.

3. Exploratory Data Analysis Summary

3.1 Dataset Overview

I began my analysis by examining the dataset structure and quality:

```
# Initial data overview
print("INITIAL DATA OVERVIEW:")
print(f"Dataset shape: {df.shape}")
print(f"Columns: {list(df.columns)}")
print("\nFirst 5 rows:")
print(df.head())
print("\nMissing values:")
print(df.isnull().sum())
```

My Results:

```

INITIAL DATA OVERVIEW:
Dataset shape: (3276, 10)
Columns: ['ph', 'Hardness', 'Solids', 'Chloramines', 'Sulfate', 'Conductivity', 'Organic_carbon', 'Trihalomethanes', 'Turbidity', 'Potability']

First 5 rows:
   ph      Hardness      Solids  Chloramines      Sulfate  Conductivity \
0   NaN    204.890455    20791.318981    7.300212    368.516441    564.308654
1  3.716080    129.422921    18630.057858    6.635246         NaN    592.885359
2  8.099124    224.236259    19909.541732    9.275884         NaN    418.606213
3  8.316766    214.373394    22018.417441    8.059332    356.886136    363.266516
4  9.092223    181.101509    17978.986339    6.546600    310.135738    398.410813

   Organic_carbon  Trihalomethanes  Turbidity  Potability
0      10.379783         86.990970    2.963135         0
1      15.180013         56.329076    4.500656         0
2      16.868637         66.420093    3.055934         0
3      18.436524        100.341674    4.628771         0
4      11.558279         31.997993    4.075075         0

Missing values:
ph           491
Hardness      0
Solids        0
Chloramines   0
Sulfate       781
Conductivity  0
Organic_carbon 0
Trihalomethanes 162
Turbidity     0
Potability    0
dtype: int64

```

I discovered that three features had missing values that needed to be addressed during preprocessing.

3.2 Target Variable Analysis

I analyzed the target variable distribution to understand the classification problem:

```

# Basic EDA for understanding
print("\nBASIC EDA:")
print("Target variable distribution:")
print(df['Potability'].value_counts())
print(f"Class balance: {df['Potability'].value_counts(normalize=True)}")

```

My Results:

```

BASIC EDA:
Target variable distribution:
Potability
0      1998
1      1278
Name: count, dtype: int64
Class balance: Potability
0      0.60989
1      0.39011
Name: proportion, dtype: float64

```

I found the dataset has a moderate class imbalance with approximately 61% non-potable and 39% potable water samples.

3.3 Feature Correlation Analysis

I investigated the relationships between predictor variables and the target:

```
# Correlation analysis for feature engineering insights
correlation_matrix = df.corr()
print("\nTop correlations with target:")
target_corr = abs(correlation_matrix['Potability']).sort_values(ascending=False)
print(target_corr.head(6))
```

My Results:

```
Top correlations with target:
Potability      1.000000
Solids           0.033743
Organic_carbon  0.030001
Chloramines      0.023779
Sulfate          0.020476
Hardness         0.013837
Name: Potability, dtype: float64
```

I observed that all features show weak linear correlations with the target variable, suggesting that water potability determination involves complex, potentially non-linear interactions between multiple physicochemical parameters rather than simple dependence on individual variables.

4. Class Labelling for Target Variable

4.1 My Target Variable Classification Strategy

I assessed the target variable to determine the optimal classification approach:

```
# Check if target variable is numerical or categorical
print(f"Target variable 'Potability' type: {df['Potability'].dtype}")
print(f"Unique values: {df['Potability'].unique()}")
print(f"Current distribution: {df['Potability'].value_counts().sort_index().to_dict()}")
```

My Results:

```
Target variable 'Potability' type: int64
Unique values: [0 1]
Current distribution: {0: 1998, 1: 1278}
```

My Decision: I found that the target variable is already categorical with a binary classification structure. Since the minority class represents 39% of observations (above the 0.30 threshold for acceptable balance), I decided to preserve the original binary classification scheme to maintain engineering relevance and regulatory compliance interpretation.

My Alternative Demonstration: To show my understanding of multi-class labelling, I created a 4-class pH-based categorization:

```
df['pH_quartiles'] = pd.qcut(df['ph'], q=4, labels=['Very_Acidic', 'Acidic', 'Basic', 'Very_Basic'])
df['pH_4class'] = pd.Categorical(df['pH_quartiles']).codes

print("pH-based 4-class distribution:")
ph_class_dist = df['pH_4class'].value_counts().sort_index()
print(ph_class_dist.to_dict())
print(f"Balanced distribution: {df['pH_4class'].value_counts(normalize=True).round(3).to_dict()}")
```

My Results:

```
pH-based 4-class distribution:
{0: 819, 1: 1065, 2: 573, 3: 819}
```

5. Feature Engineering and Feature Selection

5.1 My Data Preprocessing and Normalization

5.1.1 Missing Value Treatment

I handled missing values using median imputation:

```
numerical_cols = df.select_dtypes(include=[np.number]).columns
for col in numerical_cols:
    if df[col].isnull().sum() > 0:
        df[col].fillna(df[col].median(), inplace=True)
        print(f"Filled {col} missing values with median")
```

My Results:

```
DATA CLEANING:
Filled ph missing values with median
Filled Sulfate missing values with median
Filled Trihalomethanes missing values with median
Missing values after cleaning: 0
```

5.1.2 Feature Normalization

I applied StandardScaler normalization to ensure fair feature contribution:

```
# Normalization of numerical features
scaler = StandardScaler()
exclude_from_scaling = ['Potability', 'target_variable', 'turbidity_category', 'pH_quartiles',
                        'hardness_category', 'turbidity_cat_num', 'hardness_cat_num',
                        'ph_category', 'ph_cat_num', 'chloramines_category', 'chloramines_cat_num']
feature_cols = [col for col in df.columns if col not in exclude_from_scaling]
df_scaled = df.copy()
df_scaled[feature_cols] = scaler.fit_transform(df[feature_cols])
```

My Results:

```
Features normalized using StandardScaler:
Normalized features: ['ph', 'Hardness', 'Solids', 'Chloramines', 'Sulfate', 'Conductivity', 'Organic_carbon', 'Trihalomethanes', 'Turbidity', 'pH_4class']
```

5.2 My Integer Categorization of Features

I systematically converted continuous variables to categorical representations:

```
# Categorize Turbidity into 3 levels
df['turbidity_category'] = pd.cut(df['Turbidity'], bins=3, labels=['Low', 'Medium', 'High'])
df['turbidity_cat_num'] = pd.Categorical(df['turbidity_category']).codes

# Categorize Hardness into 4 levels
df['hardness_category'] = pd.cut(df['Hardness'], bins=4, labels=['Soft', 'Moderate', 'Hard', 'Very_Hard'])
df['hardness_cat_num'] = pd.Categorical(df['hardness_category']).codes

# Categorize pH into acid/neutral/basic levels
df['ph_category'] = pd.cut(df['ph'], bins=[0, 6.5, 7.5, 14], labels=['Acidic', 'Neutral', 'Basic'])
df['ph_cat_num'] = pd.Categorical(df['ph_category']).codes

# Categorize Chloramines levels
df['chloramines_category'] = pd.cut(df['Chloramines'], bins=3, labels=['Low', 'Medium', 'High'])
df['chloramines_cat_num'] = pd.Categorical(df['chloramines_category']).codes
```

My Results:

```
Integer categorization completed:
- Turbidity categories: turbidity_cat_num
0      554
1     2392
2      330
Name: count, dtype: int64
- Hardness categories: hardness_cat_num
0       41
1     1100
2     2001
3      134
Name: count, dtype: int64
- pH categories: {-1: 1, 0: 967, 1: 1249, 2: 1059}
- Chloramines categories: {0: 181, 1: 2671, 2: 424}
```

5.3 My Composite Feature Development

5.3.1 Ratio-Based Features

I created ratio-based features inspired by my understanding of water chemistry:

```
# Ratio features (based on correlation analysis)
df_scaled['ph_hardness_ratio'] = df_scaled['ph'] / (df_scaled['Hardness'] + 0.001)
df_scaled['solids_conductivity_ratio'] = df_scaled['Solids'] / (df_scaled['Conductivity'] + 0.001)
df_scaled['chloramines_sulfate_ratio'] = df_scaled['Chloramines'] / (df_scaled['Sulfate'] + 0.001)
```

5.3.2 Interaction Features

I developed interaction terms to capture synergistic effects:

```
# Product features (interaction effects)
df_scaled['organic_carbon_trihalomethanes_product'] = df_scaled['Organic_carbon'] * df_scaled['Trihalomethanes']
df_scaled['ph_conductivity_interaction'] = df_scaled['ph'] * df_scaled['Conductivity']
```

My Final Engineered Features:

1. ph_hardness_ratio
2. solids_conductivity_ratio
3. chloramines_sulfate_ratio
4. organic_carbon_trihalomethanes_product

5. ph_conductivity_interaction
6. turbidity_cat_num
7. hardness_cat_num
8. ph_cat_num
9. chloramines_cat_num

6. Decision Tree Model Development and Training

6.1 Model Architecture and Hyperparameter Configuration

A consistent decision tree architecture was implemented across all feature set evaluations to ensure fair comparative analysis:

```
# Train decision tree
dt = DecisionTreeClassifier(
    random_state=42,
    max_depth=10,
    min_samples_split=10,
    min_samples_leaf=5
)
```

Hyperparameter Justification:

- **max_depth=10:** Balances model complexity with interpretability
- **min_samples_split=10:** Prevents excessive partitioning of small subsets
- **min_samples_leaf=5:** Ensures statistical significance of leaf nodes

6.2 Feature Set Design and Experimental Framework

Six distinct feature sets were systematically designed to evaluate different aspects of feature engineering effectiveness:

```
# Updated feature_sets with more comprehensive sets
feature_sets = {
    'Set1_Original_9Features': original_features,
    'Set2_Top5Correlated': top_corr_features,
    'Set3_Original_Plus_Ratios': original_features + ['ph_hardness_ratio', 'solids_conductivity_ratio', 'chloramines_sulfate_ratio'],
    'Set4_Engineered_Features': ['ph_hardness_ratio', 'organic_carbon_trihalomethanes_product', 'turbidity_cat_num', 'hardness_cat_num', 'ph_conductivity_interaction'],
    'Set5_Best_Mixed': ['Sulfate', 'Conductivity', 'Organic_carbon', 'ph_hardness_ratio', 'organic_carbon_trihalomethanes_product', 'turbidity_cat_num'],
    'Set6_All_Categorical': ['turbidity_cat_num', 'hardness_cat_num', 'ph_cat_num', 'chloramines_cat_num'] # NEW SET
}
```

Experimental Design:

1. **Set1_Original_9Features:** Baseline performance using raw physicochemical parameters
2. **Set2_Top5Correlated:** Feature selection based on correlation magnitude
3. **Set3_Original_Plus_Ratios:** Integration of original and ratio-based features
4. **Set4_Engineered_Features:** Exclusive use of composite features
5. **Set5_Best_Mixed:** Curated combination of high-performing features

6. Set6_All_Categorical: Discretised categorical representations

6.3 Model Training and Validation Protocol

```
# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X_subset, y, test_size=0.3, random_state=42, stratify=y
)
```

```
dt.fit(X_train, y_train)

# Make predictions
y_pred = dt.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
```

The stratified sampling approach ensures representative class distribution in both training and testing subsets, critical for unbiased performance estimation.

7. Comparative Performance Analysis

7.1 Quantitative Results Summary

Feature Set	Features (n)	Accuracy	Tree Depth	Leaves	Dominant Feature	Importance
Set3_Original_Plus_Ratios	12	0.6419	10	114	Sulfate	0.1623
Set1_Original_9Features	9	0.6338	10	125	Sulfate	0.1987
Set6_All_Categorical	4	0.6073	9	57	ph_cat_num	0.4643
Set5_Best_Mixed	6	0.5972	10	94	ph_hardness_ratio	0.2986
Set4_Engineered_Features	5	0.5921	10	95	ph_hardness_ratio	0.3293
Set2_Top5Correlated	5	0.5788	10	108	ph_hardness_ratio	0.3317

8. Summary of My Observations

Through my experimental analysis, I discovered several critical insights about feature engineering effectiveness in water potability classification:

My Key Findings:

- Feature Integration Works Best:** My Set3_Original_Plus_Ratios achieved the highest performance (64.19% accuracy), which taught me that combining domain-specific engineered features with original measurements enhances predictive capability. I achieved a 0.81 percentage point improvement over baseline features.
- Categorical Features Are Surprisingly Effective:** My Set6_All_Categorical achieved competitive performance (60.73% accuracy) using only 4 discretized features. This showed me that decision trees can effectively exploit categorical boundaries while reducing computational complexity.
- Correlation-Based Selection Has Limitations:** My Set2_Top5Correlated produced the lowest performance (57.88%), which taught me that simple correlation-based feature selection doesn't adequately capture the complex multivariate relationships essential for water potability determination.

4. Chemical Knowledge Matters: I consistently found Sulfate as the dominant decision factor across multiple feature sets, which aligns with what I learned about water quality standards in my coursework, validating that the model captures chemically relevant patterns.
5. Balance Between Complexity and Performance: My optimal model balanced predictive performance with reasonable interpretability (10 depth, 114 leaves), which I believe is crucial for engineering applications where I need to explain my decisions.

9. Appendix: Source Code Repository

My Source Code Access: I have made my complete Jupyter notebook implementation and associated data files available through the following shared repository:

https://github.com/S1zzX/COS40007-Artificial-Intelligence-for-Engineering-/tree/main/Portfolio_Assessment/Assessment1